

Intro to Neural Networks

Elliott Ash

Text Data Course, Bocconi 2018

“Neural Networks”

- ▶ “Neural”:
 - ▶ nothing like brains
- ▶ “Networks”:
 - ▶ nothing to do with “networks” as normally understood – in particular, nothing to do with network theory in social science.


Recent History

- ▶ NNs frequently outperform other ML techniques on very large and complex problems.
- ▶ Increase in computing power makes them computationally tractable, graphical processing units (GPUs, designed for video games) give you over 100x performance gain over CPUs.
- ▶ Training algorithms have improved – small tweaks have made a huge impact.
- ▶ Some theoretical limitations of ANNs have turned out to be benign in practice – for example, they work well on non-convex functions.

Will it last?

- ▶ Three key principles of deep learning will persist:
 - ▶ **Simplicity**
 - ▶ feature engineering is obsolete
 - ▶ complex, brittle, engineering-heavy pipelines replaced with simple, end-to-end trainable models, composed of 5-6 tensor operations.
 - ▶ **Scalability**
 - ▶ amenable to parallelization on GPUs or TPUs (tensor processing units)
 - ▶ trained on batches of data, so can be scaled to datasets of arbitrary size.
 - ▶ **Versatility and reusability**
 - ▶ can be trained on additional data without restarting from scratch, therefore amenable for continuous online learning.
 - ▶ deep-learning models are repurposable and thus reusable

Perceptron LTU



- In a perceptron, an individual neuron (called an LTU, or linear threshold unit) is defined by


$$h(\mathbf{x}) = \text{step}(\boldsymbol{\omega}' \mathbf{x})$$

where $\text{step}(\cdot)$ is the step function.

- The neuron computes a linear combination of the inputs; if result exceeds threshold, output positive class, otherwise negative class.
- A model with a single LTU is similar to a logistic regression model.

Perceptron

- ▶ A perceptron is an array of LTUs in parallel:





- ▶ Can be trained by reinforcing connections that lead to the correct output:

$$\omega_{ij}^{t+1} = \omega_{ij}^t + \eta(\hat{y}_j - y_j)x_i$$

- ▶ ω_{jt}^t , weight between the i th input neuron and j th output neuron, at learning stage t
- ▶ x_i , input i for this row
- ▶ $\hat{y}_j - y_j$, predicted output minus actual output.
- ▶ η , learning rate



Multi-Layer Perceptrons → “Deep Learning”

- ▶ The predictive performance of perceptrons improved substantially by stacking them into multiple layers:




- ▶ Input variables are connected to multiple neurons in the hidden layer(s), which in turn are connected to the output layer.
 - ▶ This is the “deep” in deep learning!

Constructing the Last Layer



- ▶ MLPs will output a probability distribution across output classes.
 - ▶ can also output a real number, which would make a regression model.

Modern MLPs: New activation functions



- ▶ logistic function: $\sigma(z) = \frac{1}{1+\exp(-z)}$
- ▶ hyperbolic tangent function: $\tanh(z) = 2\sigma(2z) - 1$
 - ▶ ranges between -1 and 1 (rather than between 0 and 1, as the case with the logistic)
 - ▶ centered on zero, can speed up convergence
- ▶ ReLU (rectified linear unit) function: $\max\{0, z\}$,
 - ▶ deceptively simple, fast to compute, and very effective in practice
 - ▶ gradient does not saturate to zero for large values (but is flat below zero)

Google Developers Advice: MLP baseline for Text Classification

1. Calculate the number of samples/number of words per sample ratio.
2. If this ratio is less than 1500, tokenize the text as n-grams and use a simple multi-layer perceptron (MLP) model to classify them.
 - ▶ In the case of N-grams models, Google testers found that MLPs tended to out-perform logistic regression and gradient boosting machines.

Setting up a model

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential() # create a sequential model
model.add(Dense(50, # neurons in first layer
                input_dim=X.shape[1], # number of predictors
                activation='relu',)) # activation function
model.add(Dense(50, activation='relu')) # hidden layer
model.add(Dense(1, activation='sigmoid')) # output layer
model.summary()
```

► Output layer:

- for binary classification, use `activation='sigmoid'`
- for regression, do not use an activation function
- for multi-class classification, use `activation=softmax'`

Visualize a model

```
from IPython.display import SVG
from keras.utils.vis_utils import model_to_dot
dot = model_to_dot(create_reg_model(),
                   show_shapes=True,
                   show_layer_names=False)
SVG(dot.create(prog='dot', format='svg'))
```

Compile the model

```
model.compile(loss='binary_crossentropy', # cost function  
              optimizer='adam', # use adam as the optimizer  
              metrics=['accuracy']) # compute accuracy
```

- ▶ Loss function:
 - ▶ for binary classification, use `binary_crossentropy`
 - ▶ for regression, use `mean_squared_error`
 - ▶ for multi-class classification, use `sparse_categorical_crossentropy`
- ▶ Optimizer:
 - ▶ use `adam`
- ▶ Metrics:
 - ▶ for classification, use `accuracy`
 - ▶ for regression, you have to define a custom metric (see accompanying code)

Fit a model

```
model.fit(X, Y,  
          epochs=5,  
          validation_split=.2)  
  
model.get_weights()  
  
# Plot performance by epoch  
plt.plot(model_info.epoch, model_info.history['acc'])  
plt.plot(model_info.epoch, model_info.history['val_acc'])  
plt.legend(['train', 'val'], loc='best')  
  
# make predictions  
ypred = model.predict(X)
```

Saving and Loading Models

Save a model

```
model.save('keras-clf.pkl')
```

load model

```
from keras.models import load_model  
model = load_model('keras-clf.pkl')
```